

PCP and Hardness of Approximation

Aman Bansal Adwait Godbole

October 2019

1 Relaxations of Hard Problems

- Approximate Problems
- Gap Problem
- Connection between Approximation and Gap

2 A New Proof System

- Probabilistically Checkable Proof Systems
- Some Preliminaries
- Proof of the PCP Theorem

3 Hardness of Approximation

1 Relaxations of Hard Problems

- Approximate Problems
- Gap Problem
- Connection between Approximation and Gap

2 A New Proof System

- Probabilistically Checkable Proof Systems
- Some Preliminaries
- Proof of the PCP Theorem

3 Hardness of Approximation

Approximate Problems

Given an instance $\mathbf{x} \in \mathcal{X}$ of an NP-hard optimization problem with objective function $Obj : \mathcal{Y} \rightarrow \mathbb{R}$ (which is to be maximized¹), a solution \mathbf{y} is said to be an α -approximate (for $\alpha \leq 1$) solution to the instance if

$$\alpha \cdot Obj(\mathbf{y}^*) \leq Obj(\mathbf{y}) \leq Obj(\mathbf{y}^*)$$

where \mathbf{y}^* is the true (not approximated) solution to the problem instance.

¹If the problem is a minimization problem then we have a slightly different definition.

Approximate Problems

α -approximate algorithms

An algorithm **A** is an α -approximate algorithm if for all \mathbf{x} in the instance space \mathcal{X} , it returns an α -approximate solution \mathbf{y} .

α -approximate problems

Problems that render (*poly-time*) α -approximate algorithms are called α -approximate problems.

1 Relaxations of Hard Problems

- Approximate Problems
- Gap Problem
- Connection between Approximation and Gap

2 A New Proof System

- Probabilistically Checkable Proof Systems
- Some Preliminaries
- Proof of the PCP Theorem

3 Hardness of Approximation

Promise Problem

A *promise* problem Π is specified by a pair of sets (YES, NO) such that $\text{YES}, \text{NO} \subseteq \mathcal{X}$ and $\text{YES} \cap \text{NO} = \emptyset$.

Any algorithm **A** solving Π , on input \mathbf{x} , should output 'yes' if $\mathbf{x} \in \text{YES}$, 'no' if $\mathbf{x} \in \text{NO}$ and any output if \mathbf{x} is a don't care instance

Recollect the Promise Π problem from the midsem.

Gap Problems

A gap problem is a promise problem parametrized by $\alpha (< 1)$. Let P be an NP-hard optimization problem with objective function $Obj : \mathcal{Y} \rightarrow \mathbb{R}$ (which is to be maximized), the corresponding gap problem $gap_\alpha\text{-}P$ is a promise problem with (YES, NO) sets as given below:

$$YES = \{ \langle \mathbf{x}, k \rangle \mid \exists \mathbf{y} \in \mathcal{Y} \text{ such that } Obj(\mathbf{y}) \geq k \}$$

$$NO = \{ \langle \mathbf{x}, k \rangle \mid \forall \mathbf{y} \in \mathcal{Y}, Obj(\mathbf{y}) < \alpha k \}$$

Intuitively, the 'gap' refers to the interval $(\alpha k, k)$.

1 Relaxations of Hard Problems

- Approximate Problems
- Gap Problem
- Connection between Approximation and Gap

2 A New Proof System

- Probabilistically Checkable Proof Systems
- Some Preliminaries
- Proof of the PCP Theorem

3 Hardness of Approximation

Intuitively, it seems that approximate problems and gap problems are similar.

α -approximation is a relaxed variant of a search problem

gap_α is a relaxed variant of a decision problem

Indeed, this notion can be formalized.

Connection between α -approximation and gap_α

For any problem P and $0 < \alpha < 1$, α -approximating P is at least as hard as solving gap_α - P .

Proof:

Let A be an α -approximate algorithm for P . The following algorithm solves gap_α - P :

On input (ϕ, k) :

1. Let $k' = A(\phi)$
2. Accept iff $k' \geq \alpha k$

Trajectory

- 1 Relaxations of Hard Problems
 - Approximate Problems
 - Gap Problem
 - Connection between Approximation and Gap
- 2 A New Proof System
 - Probabilistically Checkable Proof Systems
 - Some Preliminaries
 - Proof of the PCP Theorem
- 3 Hardness of Approximation

1 Relaxations of Hard Problems

- Approximate Problems
- Gap Problem
- Connection between Approximation and Gap

2 A New Proof System

- Probabilistically Checkable Proof Systems
- Some Preliminaries
- Proof of the PCP Theorem

3 Hardness of Approximation

Probabilistically Checkable Proof System

(r, q, m, t) -restricted Verifier

Let $r, q, m, t : \mathbb{N} \rightarrow \mathbb{N}$. A language $L \in PCP_{c,s}[r, q, m, t]$ if L has an (r, q, m, t) restricted verifier V such that

$$\forall x \in L, \exists \pi \text{ of size at most } m(|x|), Pr_R[V^\pi[x; R] = acc] \geq c(|x|)$$

$$\forall x \notin L, \forall \pi \text{ of size at most } m(|x|), Pr_R[V^\pi[x; R] = acc] < s(|x|)$$

Resource bounds

1. $r(|x|)$ is a bound on randomness used by V
2. $q(|x|)$ is a bound on the number of locations queried by V
3. $m(|x|)$ is a bound the length of the proof to which V has oracle access
4. $t(|x|)$ is a bound on the runtime of V

Completeness and Soundness constraints

$c(|x|)$ and $s(|x|)$ are the completeness and soundness specifications

- The verifier could be adaptive or non-adaptive
- If the verifier is non-adaptive then $m(n) \leq q(n) \cdot 2^{r(n)}$
- $q(n) \leq t(n)$
- $PCP_{c,s}[r, q] \subseteq NTIME(q(n) \cdot 2^{r(n)})$
- $NP = PCP_{1,0}[0, poly(n)]$
- $BPP = PCP_{\frac{2}{3}, \frac{1}{3}}[poly(n), 0]$

PCP Theorem

$$NP = PCP_{1, \frac{1}{2}}[\log(n), 1]$$

PCP Theorem

PCP Theorem

$$NP = PCP_{1, \frac{1}{2}}[\log(n), 1]$$

We will today present a weaker version of the PCP theorem.

PCP Theorem [Weaker]

$$NP = PCP_{1, \frac{1}{2}}[\text{poly}(n), 1]$$

1 Relaxations of Hard Problems

- Approximate Problems
- Gap Problem
- Connection between Approximation and Gap

2 A New Proof System

- Probabilistically Checkable Proof Systems
- Some Preliminaries
- Proof of the PCP Theorem

3 Hardness of Approximation

Subset XOR

Consider the function $f_{\mathbf{u}}(\mathbf{x}) = \mathbf{x} \odot \mathbf{u}$.

Here for $\mathbf{x}, \mathbf{y} \in \{0, 1\}^n$, $\mathbf{x} \odot \mathbf{y} = \sum_i x_i y_i \pmod{2}$.

Subset XOR

Consider the function $f_{\mathbf{u}}(\mathbf{x}) = \mathbf{x} \odot \mathbf{u}$.

Here for $\mathbf{x}, \mathbf{y} \in \{0, 1\}^n$, $\mathbf{x} \odot \mathbf{y} = \sum_i x_i y_i \pmod{2}$.

Note that $f_{\mathbf{u}}$ is equivalent to choosing a subset (given by \mathbf{u}) of elements from $[n]$ and evaluating parity over this subset.

Subset XOR

Consider the function $f_{\mathbf{u}}(\mathbf{x}) = \mathbf{x} \odot \mathbf{u}$.

Here for $\mathbf{x}, \mathbf{y} \in \{0, 1\}^n$, $\mathbf{x} \odot \mathbf{y} = \sum_i x_i y_i \pmod{2}$.

Note that $f_{\mathbf{u}}$ is equivalent to choosing a subset (given by \mathbf{u}) of elements from $[n]$ and evaluating parity over this subset.

Random Subsum Principle

For $\mathbf{x}, \mathbf{y} \in \{0, 1\}^n$ with $\mathbf{y} \neq 0^n$, $\Pr_{\mathbf{x} \in \{0, 1\}^n}[\mathbf{x} \odot \mathbf{y} = 1] = \frac{1}{2}$

Walsh-Hadamard Code

Main Idea: Bit strings $\mathbf{u} \in \{0, 1\}^n$ are encoded as the truth table of a *linear* function over \mathbf{F}_2

Walsh-Hadamard Code

Main Idea: Bit strings $\mathbf{u} \in \{0, 1\}^n$ are encoded as the truth table of a *linear* function over \mathbf{F}_2

WH encoding

For an n -bit string $\mathbf{u} \in \{0, 1\}^n$, $\text{WH}(\mathbf{u})$ is the 2^n bit string representing the truth table of the function $f(\mathbf{x}) = \mathbf{x} \odot \mathbf{u}$ for $\mathbf{x} \in \{0, 1\}^n$.

Walsh-Hadamard Code

Main Idea: Bit strings $\mathbf{u} \in \{0, 1\}^n$ are encoded as the truth table of a *linear* function over \mathbf{F}_2

WH encoding

For an n -bit string $\mathbf{u} \in \{0, 1\}^n$, $\text{WH}(\mathbf{u})$ is the 2^n bit string representing the truth table of the function $f(\mathbf{x}) = \mathbf{x} \odot \mathbf{u}$ for $\mathbf{x} \in \{0, 1\}^n$.

Walsh-Hadamard codeword

$f \in \{0, 1\}^{2^n}$ such that $f = \text{WH}(\mathbf{u})$ for some $\mathbf{u} \in \{0, 1\}^n$.

Properties of WH

EC *Error correcting* with minimum distance $\frac{1}{2}$.

This means that for $\mathbf{x}, \mathbf{y} \in_R \{0, 1\}^n$ with $\mathbf{x} \neq \mathbf{y}$, $\text{WH}(\mathbf{x})$ and $\text{WH}(\mathbf{y})$ differ in $1/2$ the bits.

Properties of WH

EC *Error correcting* with minimum distance $\frac{1}{2}$.

This means that for $\mathbf{x}, \mathbf{y} \in_R \{0, 1\}^n$ with $\mathbf{x} \neq \mathbf{y}$, $\text{WH}(\mathbf{x})$ and $\text{WH}(\mathbf{y})$ differ in $1/2$ the bits.

LIN *Linearity* of $\text{WH}(\mathbf{u})$

$f = \text{WH}(\mathbf{u})$ when viewed as a function from $\{0, 1\}^n$ to $\{0, 1\}$ is in fact *linear* (over \mathbf{F}_2)

Properties of WH

EC *Error correcting* with minimum distance $\frac{1}{2}$.

This means that for $\mathbf{x}, \mathbf{y} \in_R \{0, 1\}^n$ with $\mathbf{x} \neq \mathbf{y}$, $\text{WH}(\mathbf{x})$ and $\text{WH}(\mathbf{y})$ differ in $1/2$ the bits.

LIN *Linearity* of $\text{WH}(\mathbf{u})$

$f = \text{WH}(\mathbf{u})$ when viewed as a function from $\{0, 1\}^n$ to $\{0, 1\}$ is in fact *linear* (over \mathbf{F}_2)

LT *Locally Testable*

Given access to a function $f : \{0, 1\}^n \rightarrow \{0, 1\}$, we can check whether it is a *Walsh-Hadamard* code-word by querying a *constant* number of places.

Properties of WH

EC *Error correcting* with minimum distance $\frac{1}{2}$.

This means that for $\mathbf{x}, \mathbf{y} \in_R \{0, 1\}^n$ with $\mathbf{x} \neq \mathbf{y}$, $\text{WH}(\mathbf{x})$ and $\text{WH}(\mathbf{y})$ differ in $1/2$ the bits.

LIN *Linearity* of $\text{WH}(\mathbf{u})$

$f = \text{WH}(\mathbf{u})$ when viewed as a function from $\{0, 1\}^n$ to $\{0, 1\}$ is in fact *linear* (over \mathbf{F}_2)

LT *Locally Testable*

Given access to a function $f : \{0, 1\}^n \rightarrow \{0, 1\}$, we can check whether it is a *Walsh-Hadamard* code-word by querying a *constant* number of places.

LD *Locally Decodable*

Given f and an $\mathbf{x} \in \{0, 1\}^n$, we can find $\tilde{f}(\mathbf{x})$ in constant queries to f , where \tilde{f} is the true codeword.

LT: Local Testability

LIN: $\text{WH}(\mathbf{u})$ for $\mathbf{u} \in \{0, 1\}^n$ captures *all* n -bit linear functions on \mathbf{F}_2 .

LIN: $\text{WH}(\mathbf{u})$ for $\mathbf{u} \in \{0, 1\}^n$ captures *all* n -bit linear functions on \mathbf{F}_2 .

ρ -closeness of functions

Functions f, g are ρ -close if $\Pr_{\mathbf{x} \in_R \{0,1\}^n} [f(\mathbf{x}) = g(\mathbf{x})] \geq \rho$.

A function f is ρ -close to a linear function if there exists a linear function g such that f and g are ρ -close

LT: Local Testability

LIN: $\text{WH}(\mathbf{u})$ for $\mathbf{u} \in \{0, 1\}^n$ captures *all* n -bit linear functions on \mathbf{F}_2 .

ρ -closeness of functions

Functions f, g are ρ -close if $\Pr_{\mathbf{x} \in_R \{0,1\}^n} [f(\mathbf{x}) = g(\mathbf{x})] \geq \rho$.

A function f is ρ -close to a linear function if there exists a linear function g such that f and g are ρ -close

Let f be such that

$$\Pr_{\mathbf{x}, \mathbf{y} \in_R \{0,1\}^n} [f(\mathbf{x} + \mathbf{y}) = f(\mathbf{x}) + f(\mathbf{y})] \geq \rho$$

for some $\rho > \frac{1}{2}$. Then f is ρ -close to a linear function.

EC: For $\mathbf{x} \neq \mathbf{y}$, $\text{WH}(\mathbf{x})$ and $\text{WH}(\mathbf{y})$ differ in $1/2$ the bits

LD: Local Decodability

EC: For $\mathbf{x} \neq \mathbf{y}$, $\text{WH}(\mathbf{x})$ and $\text{WH}(\mathbf{y})$ differ in $1/2$ the bits

Let f be $(1 - \delta)$ -close to a linear function \tilde{f} for some $\delta < 1/4$. Then by EC, f uniquely determines \tilde{f} .

EC: For $\mathbf{x} \neq \mathbf{y}$, $\text{WH}(\mathbf{x})$ and $\text{WH}(\mathbf{y})$ differ in $1/2$ the bits

Let f be $(1 - \delta)$ -close to a linear function \tilde{f} for some $\delta < 1/4$. Then by EC, f uniquely determines \tilde{f} .

So, given a (possibly illegal) f , having a corresponding \tilde{f} , we want to find $\tilde{f}(\mathbf{x})$. Here we have oracle access only to f . The idea is to once again use randomness.

LD: Local Decodability

Objective: With oracle access only to f , given an $\mathbf{x} \in \{0,1\}^n$, find $\tilde{f}(\mathbf{x})$.
The idea is to use randomness and linearity.

LD: Local Decodability

Objective: With oracle access only to f , given an $\mathbf{x} \in \{0, 1\}^n$, find $\tilde{f}(\mathbf{x})$.
The idea is to use randomness and linearity.

- Choose $\mathbf{x}' \in_R \{0, 1\}^n$
- Set $\mathbf{x}'' \leftarrow \mathbf{x} + \mathbf{x}'$
- Let $\mathbf{y}' = f(\mathbf{x}')$ and $\mathbf{y}'' = f(\mathbf{x}'')$
- Output $\mathbf{y}' + \mathbf{y}''$

LD: Local Decodability

Objective: With oracle access only to f , given an $\mathbf{x} \in \{0, 1\}^n$, find $\tilde{f}(\mathbf{x})$.
The idea is to use randomness and linearity.

- Choose $\mathbf{x}' \in_R \{0, 1\}^n$
- Set $\mathbf{x}'' \leftarrow \mathbf{x} + \mathbf{x}'$
- Let $\mathbf{y}' = f(\mathbf{x}')$ and $\mathbf{y}'' = f(\mathbf{x}'')$
- Output $\mathbf{y}' + \mathbf{y}''$

With probability atleast $1 - 2\delta$ we have $\mathbf{y}' = f(\mathbf{x}')$ and $\mathbf{y}'' = f(\mathbf{x}'')$ and hence $\tilde{f}(\mathbf{x}) = \mathbf{y}' + \mathbf{y}''$.

1 Relaxations of Hard Problems

- Approximate Problems
- Gap Problem
- Connection between Approximation and Gap

2 A New Proof System

- Probabilistically Checkable Proof Systems
- Some Preliminaries
- Proof of the PCP Theorem

3 Hardness of Approximation

Proof of the PCP Theorem [Weaker]

PCP Theorem [Weaker]

$$NP = PCP_{1, \frac{1}{2}}[poly(n), 1]$$

Proof of the PCP Theorem [Weaker]

PCP Theorem [Weaker]

$$NP = PCP_{1, \frac{1}{2}}[poly(n), 1]$$

We show that QUADEQ - the satisfiability problem for quadratic equations over \mathbf{F}_2 - has a $PCP[poly(n), 1]$ proof system

Proof of the PCP Theorem [Weaker]

PCP Theorem [Weaker]

$$NP = PCP_{1, \frac{1}{2}}[poly(n), 1]$$

We show that QUADEQ - the satisfiability problem for quadratic equations over \mathbf{F}_2 - has a $PCP[poly(n), 1]$ proof system

QUADEQ over variables u_1, u_2, \dots, u_n is of the form $AU = b$, where A is an $m \times n^2$ matrix and $b \in \{0, 1\}^m$. $U = \mathbf{u} \otimes \mathbf{u}$ is the tensor product (or the Hadamard product).

Proof of the PCP Theorem [Weaker]

PCP Theorem [Weaker]

$$NP = PCP_{1, \frac{1}{2}}[poly(n), 1]$$

We show that QUADEQ - the satisfiability problem for quadratic equations over \mathbf{F}_2 - has a $PCP[poly(n), 1]$ proof system

QUADEQ over variables u_1, u_2, \dots, u_n is of the form $AU = b$, where A is an $m \times n^2$ matrix and $b \in \{0, 1\}^m$. $U = \mathbf{u} \otimes \mathbf{u}$ is the tensor product (or the Hadamard product).

Claim

QUADEQ, the language of all satisfiable instances is NP-complete

What is the proof π and what does the verifier \mathcal{V} do?

What is the proof π and what does the verifier \mathcal{V} do?

π The proof is $\langle \text{WH}(\mathbf{u}), \text{WH}(\mathbf{u} \otimes \mathbf{u}) \rangle$

What is the proof π and what does the verifier \mathcal{V} do?

π The proof is $\langle \text{WH}(\mathbf{u}), \text{WH}(\mathbf{u} \otimes \mathbf{u}) \rangle$

\mathcal{V} Denote the proof by $f = \text{WH}(\mathbf{u})$ and $g = \text{WH}(\mathbf{u} \otimes \mathbf{u})$.

The verifier does the following

- 1) Check linearity of f and g
- 2) Verify that g encodes $\mathbf{u} \otimes \mathbf{u}$
- 3) Verify that f encodes a satisfying assignment

Check linearity of f and g

Note that $f = \text{WH}(\mathbf{u})$ and $g = \text{WH}(\mathbf{u} \otimes \mathbf{u})$

Check linearity of f and g

Note that $f = \text{WH}(\mathbf{u})$ and $g = \text{WH}(\mathbf{u} \otimes \mathbf{u})$

\mathcal{V} performs a 0.99-close (high-probability) linearity test on both f and g . This is done by the LT property described earlier.

Check linearity of f and g

Note that $f = \text{WH}(\mathbf{u})$ and $g = \text{WH}(\mathbf{u} \otimes \mathbf{u})$

\mathcal{V} performs a 0.99-close (high-probability) linearity test on both f and g . This is done by the LT property described earlier.

Note crucially that a high but nevertheless constant closeness suffices. This is since we eventually plan to query π at only a small constant number of points.

Verify that g encodes $\mathbf{u} \otimes \mathbf{u}$

\mathcal{V} chooses \mathbf{r}, \mathbf{r}' independently at random from $\{0, 1\}^n$ and assert that $f(\mathbf{r})f(\mathbf{r}') = g(\mathbf{r} \otimes \mathbf{r}')$.

Verify that g encodes $\mathbf{u} \otimes \mathbf{u}$

\mathcal{V} chooses \mathbf{r}, \mathbf{r}' independently at random from $\{0, 1\}^n$ and assert that $f(\mathbf{r})f(\mathbf{r}') = g(\mathbf{r} \otimes \mathbf{r}')$.

Let W be an $n \times n$ matrix representing the entries of \mathbf{w} and U be such a matrix for $\mathbf{u} \otimes \mathbf{u}$. Then

1. $g(\mathbf{r} \oplus \mathbf{r}') = \mathbf{w} \odot (\mathbf{r} \otimes \mathbf{r}') = \mathbf{r}W\mathbf{r}'$
2. $f(\mathbf{r})f(\mathbf{r}') = (\mathbf{u} \odot \mathbf{r})(\mathbf{u} \odot \mathbf{r}') = \mathbf{r}U\mathbf{r}'$

Verify that g encodes $\mathbf{u} \otimes \mathbf{u}$

\mathcal{V} chooses \mathbf{r}, \mathbf{r}' independently at random from $\{0, 1\}^n$ and assert that $f(\mathbf{r})f(\mathbf{r}') = g(\mathbf{r} \otimes \mathbf{r}')$.

Let W be an $n \times n$ matrix representing the entries of \mathbf{w} and U be such a matrix for $\mathbf{u} \otimes \mathbf{u}$. Then

1. $g(\mathbf{r} \oplus \mathbf{r}') = \mathbf{w} \odot (\mathbf{r} \otimes \mathbf{r}') = \mathbf{r}W\mathbf{r}'$
2. $f(\mathbf{r})f(\mathbf{r}') = (\mathbf{u} \odot \mathbf{r})(\mathbf{u} \odot \mathbf{r}') = \mathbf{r}U\mathbf{r}'$

By the random subsum principle, we claim this test rejects atleast $1/4$ of the time on instances where $\mathbf{w} \neq \mathbf{u} \otimes \mathbf{u}$. Repeating this 3 times, we get probability of rejection as $37/64$.

Verify that f encodes a satisfying assignment

Now we are assured that the form of π is $\langle \text{WH}(\mathbf{u}), \text{WH}(\mathbf{u} \otimes \mathbf{u}) \rangle$ for some $\mathbf{u} \in \{0, 1\}^n$.

Verify that f encodes a satisfying assignment

Now we are assured that the form of π is $\langle \text{WH}(\mathbf{u}), \text{WH}(\mathbf{u} \otimes \mathbf{u}) \rangle$ for some $\mathbf{u} \in \{0, 1\}^n$.

- 1 All that remains is to check that \mathbf{u} is a satisfying assignment
- 2 Wonderfully, we also have $O(1)$ access to $A_i \cdot (\mathbf{u} \otimes \mathbf{u})$, the value of the i^{th} equation in the QUADEQ instance and can match it with b_i .

Verify that f encodes a satisfying assignment

Now we are assured that the form of π is $\langle \text{WH}(\mathbf{u}), \text{WH}(\mathbf{u} \otimes \mathbf{u}) \rangle$ for some $\mathbf{u} \in \{0, 1\}^n$.

- 1 All that remains is to check that \mathbf{u} is a satisfying assignment
- 2 Wonderfully, we also have $O(1)$ access to $A_i \cdot (\mathbf{u} \otimes \mathbf{u})$, the value of the i^{th} equation in the QUADEQ instance and can match it with b_i .

But but but ... how do we check all the m equations of the QUADEQ instance in constant number of queries? ☹

Verify that f encodes a satisfying assignment

Now we are assured that the form of π is $\langle \text{WH}(\mathbf{u}), \text{WH}(\mathbf{u} \otimes \mathbf{u}) \rangle$ for some $\mathbf{u} \in \{0, 1\}^n$.

- 1 All that remains is to check that \mathbf{u} is a satisfying assignment
- 2 Wonderfully, we also have $O(1)$ access to $A_i \cdot (\mathbf{u} \otimes \mathbf{u})$, the value of the i^{th} equation in the QUADEQ instance and can match it with b_i .

But but but ... how do we check all the m equations of the QUADEQ instance in constant number of queries? ☹️

Use the random subsum principle AGAIN! Choose a subset of equations randomly from $[k]$ and add them together to create a new quadratic equation. If \mathbf{u} did not satisfy even one equation of the original system, it will not satisfy the new equation with probability at least $1/2$.

With this, we have proved that $NP \subseteq PCP[poly(n), 1]$. The other direction is trivial.

The stronger theorem makes further observations regarding the form of the proof π given here. Then it uses further results such as gap amplification and alphabet reduction to prove the general statement $NP = PCP(\log(n), 1)$.

- 1 Relaxations of Hard Problems
 - Approximate Problems
 - Gap Problem
 - Connection between Approximation and Gap
- 2 A New Proof System
 - Probabilistically Checkable Proof Systems
 - Some Preliminaries
 - Proof of the PCP Theorem
- 3 Hardness of Approximation

Hardness of Approximation

A Sad Result

gap_{α} -MAX3SAT is NP-hard

Just as SAT captured the essence of hardness of exact solution, gap_{α} -MAX3SAT, or its more general formulation, q CSP, captures the essence of hardness of approximation

Hardness of Approximation

Proof Method

PCP Theorem \implies $\text{gap}_\alpha\text{-MAX3SAT}$ is NP-hard

Proof: Consider any $L \in \text{PCP}_{1, \frac{1}{2}}[c \cdot \log(n), Q]$. The idea is to encode the Verifier's possible actions by a Boolean formula Ψ .

$$\Psi = \bigwedge_{\text{coins } R} h_R$$

But h_r is an arbitrary predicate over Q variables.

Hardness of Approximation

Fact

$\forall q, \exists l(q), k(q)$ such that any q -ary Boolean function h can be encoded by a 3-CNF formula ψ_h with $k(q)$ clauses over $q + l(q)$ variables $x_1, \dots, x_q, z_1, \dots, z_{l(q)}$ such that

$$h(x) = 1 \implies \exists z, \psi_h(x, z) = 1$$

$$h(x) = 0 \implies \forall z, \psi_h(x, z) = 0$$

Hardness of Approximation

$$\Psi = \bigwedge_{\text{coins } R} \psi_{h_R}$$

1. If $x \in L$ then \exists proof π such that $\forall R, h_R(\pi) = 1$
2. If $x \notin L$ then at least $\frac{1}{2}$ choices of R accept make $h_R(\pi) = 0$. If the total number of clauses are $M (= 2^R k(q))$ then maximum fraction of clauses that can be satisfied is $(1 - \frac{1}{k})$. □

This proves that PCP-theorem leads to the hardness of gap_α -MAX3SAT which in turn as presented earlier leads to NP-hardness of α -approximating MAX3SAT.

- 1 Arora, Sanjeev, and Boaz Barak. Computational complexity: a modern approach. Cambridge University Press, 2009.
- 2 Limits of approximation algorithms : PCPs and Unique Games - Spring Semester (2009-10) at TIFR, IMSc
<http://www.tcs.tifr.res.in/~prahladh/teaching/2009-10/limits/>

